

# Transactional systems

COURSE 4: Databases

# Transactional systems

# Transaction

- Set of operations on the database, set of statements:
  - insert, update, delete
- Delimited by statements or function calls of type:
  - begin transaction
  - end transaction
- All operations are finalized with success or none is saved in the db.
- A transactional system must
  - manage concurrent transactions.
  - ensure consistent data in case of failure.

# Transaction

Statement 1

Statement 2

commit -- end transaction 1

Statement 3

Statement 4

Statement 5

commit -- end transaction 2

# Transaction properties

*ACID*



**ATOMICITY**

CONSISTENCY

ISOLATION

DURABILITY

- all changes or none
  - collection of steps → single indivisible unit.
- If one operation fails all changes to the database must be undone
  - Failures in transaction, example: statement error, violating unique constraint.
  - System failures, OS crashed.




ATOMICITY

**CONSISTENCY**

ISOLATION

DURABILITY

- If a transaction is run starting from a database in a consistent state, the database must be consistent at the end of the transaction.
- Database constraints
  - PRIMARY KEY key constraint, UNIQUE, NOT NULL, FOREIGN KEY referential integrity, CHECK
- Business constraints



- The database may at some point be in an *inconsistent state*.
- Inconsistencies are not visible in a database system (ensured by *atomicity*).
- The old values of any data on which a transaction performs is written to a log file used by a  
→ *recovery system*



ATOMICITY

CONSISTENCY

**ISOLATION**

DURABILITY

- The database system must ensure that transactions run without interference.
  - For any pair of transactions  $T_i, T_j$ ,  
first statement of transaction  $T_i$  is executed after  $T_j$  finished or  
first statement of transaction  $T_j$  is executed after  $T_i$  finished.



ATOMICITY

CONSISTENCY

ISOLATION

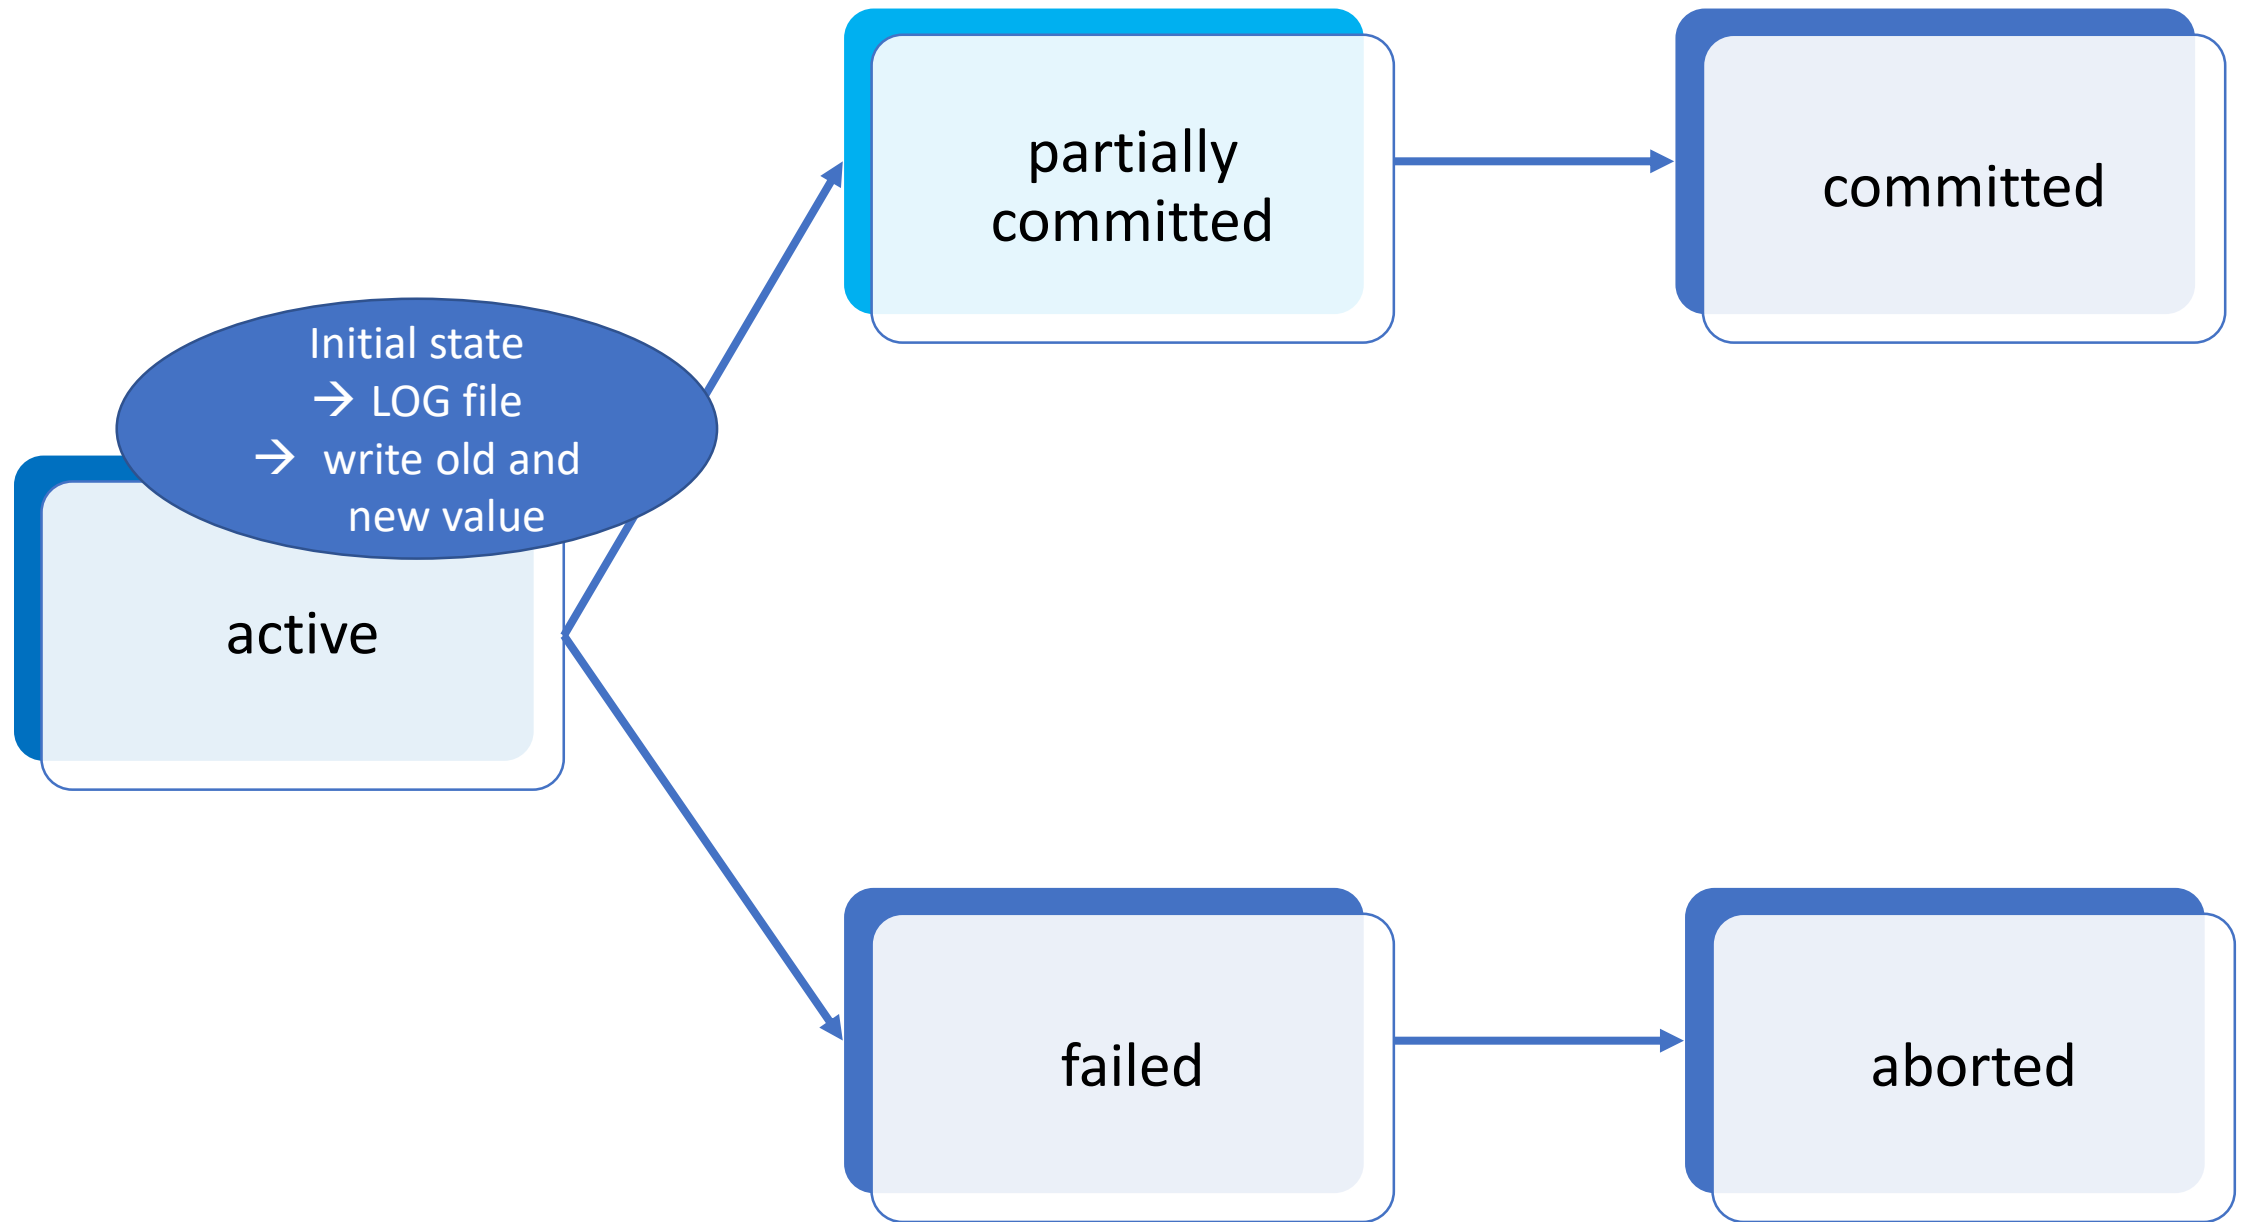
**DURABILITY**

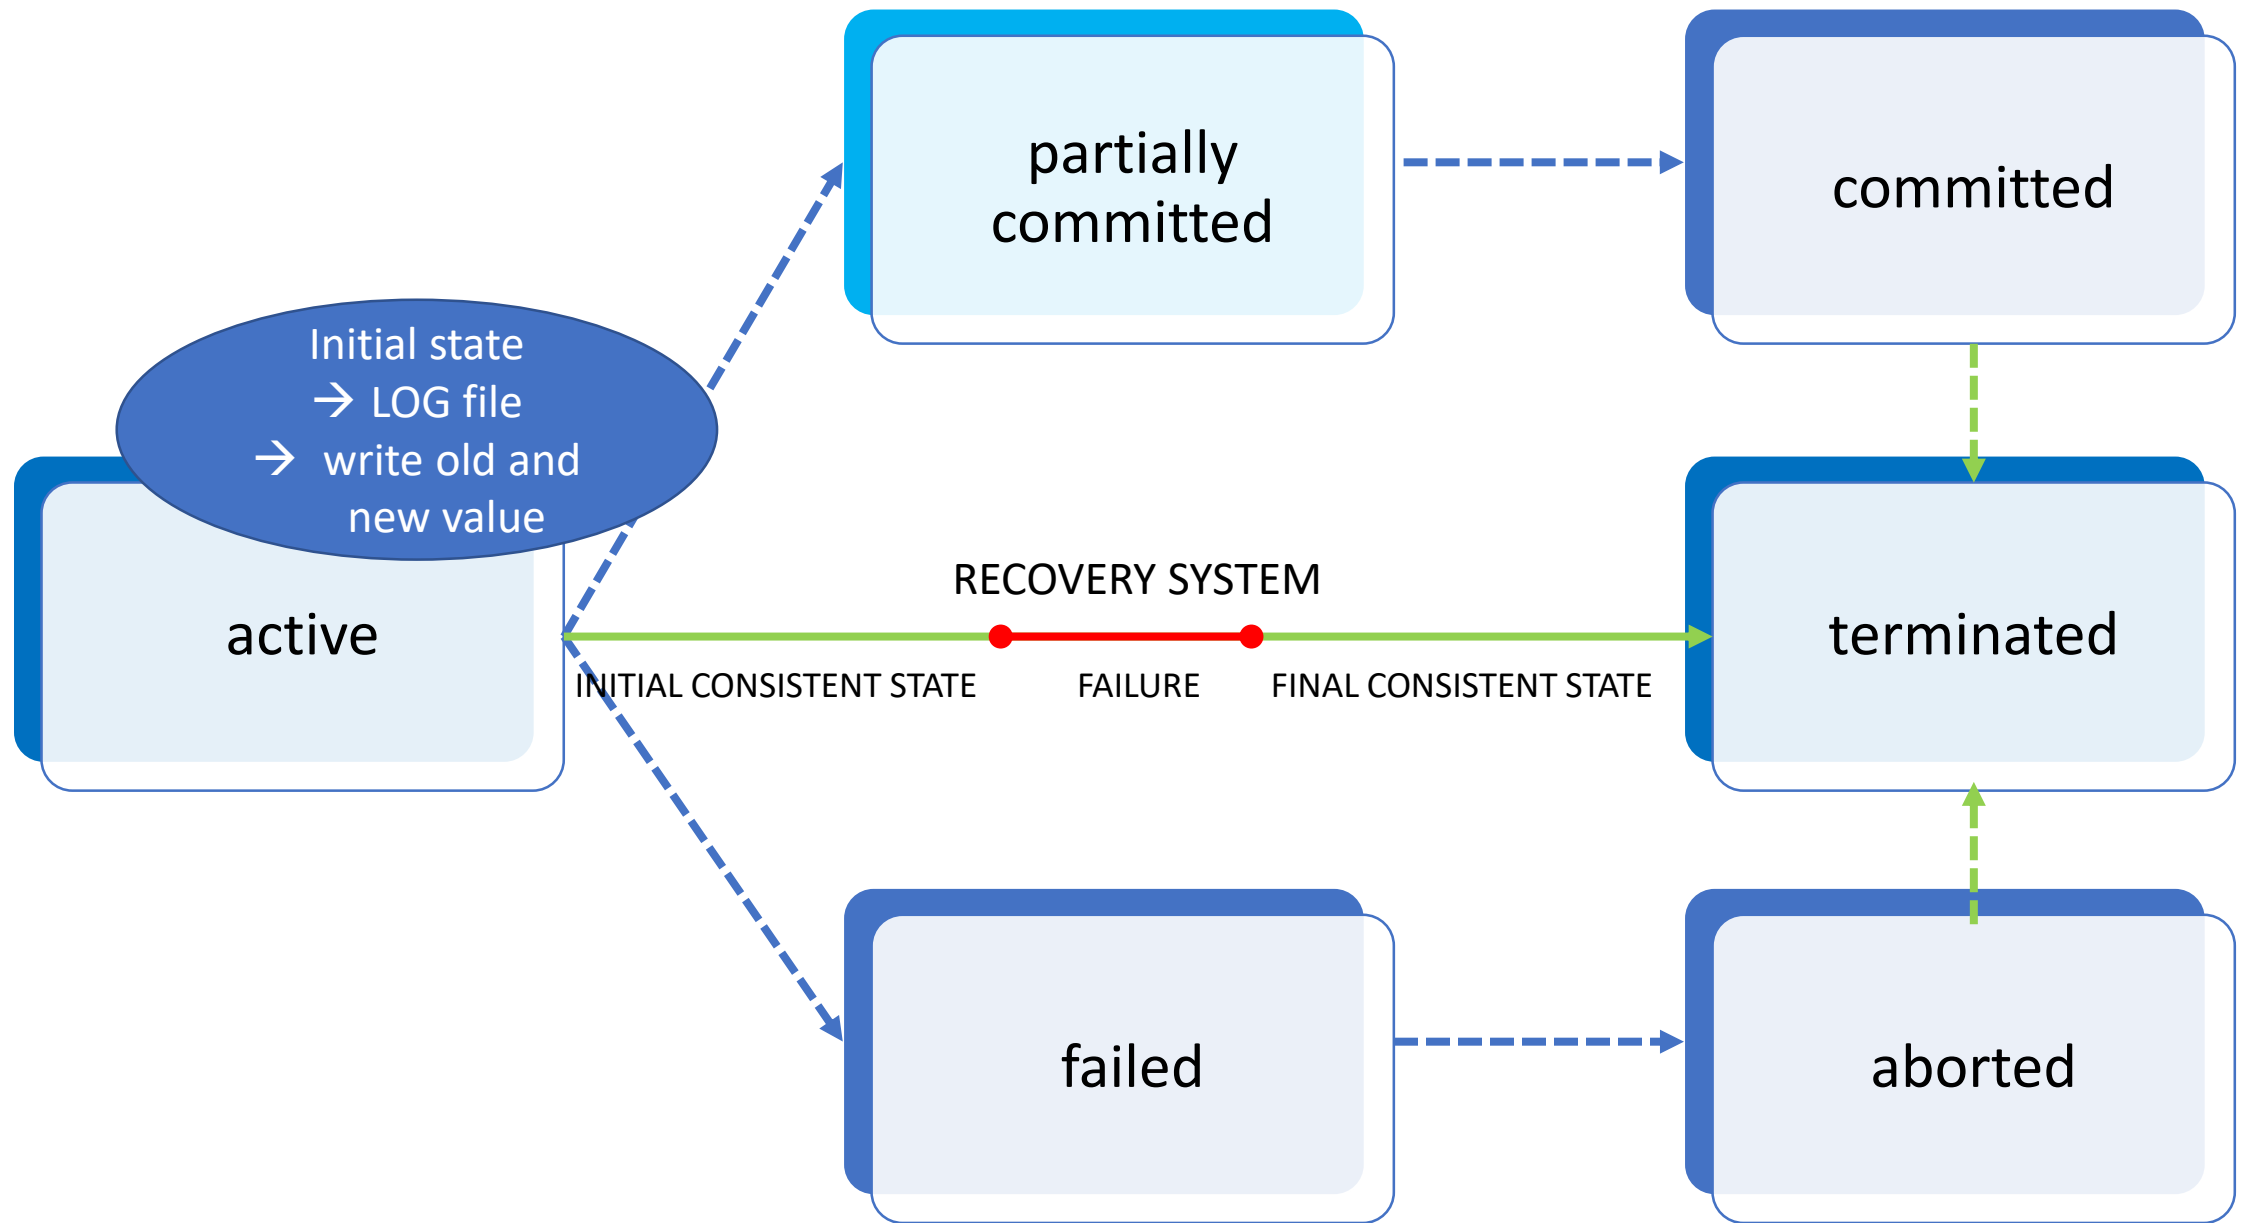
- After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.
- Information about the updates performed by the transaction is written to disk and used to reconstruct the database after failure.

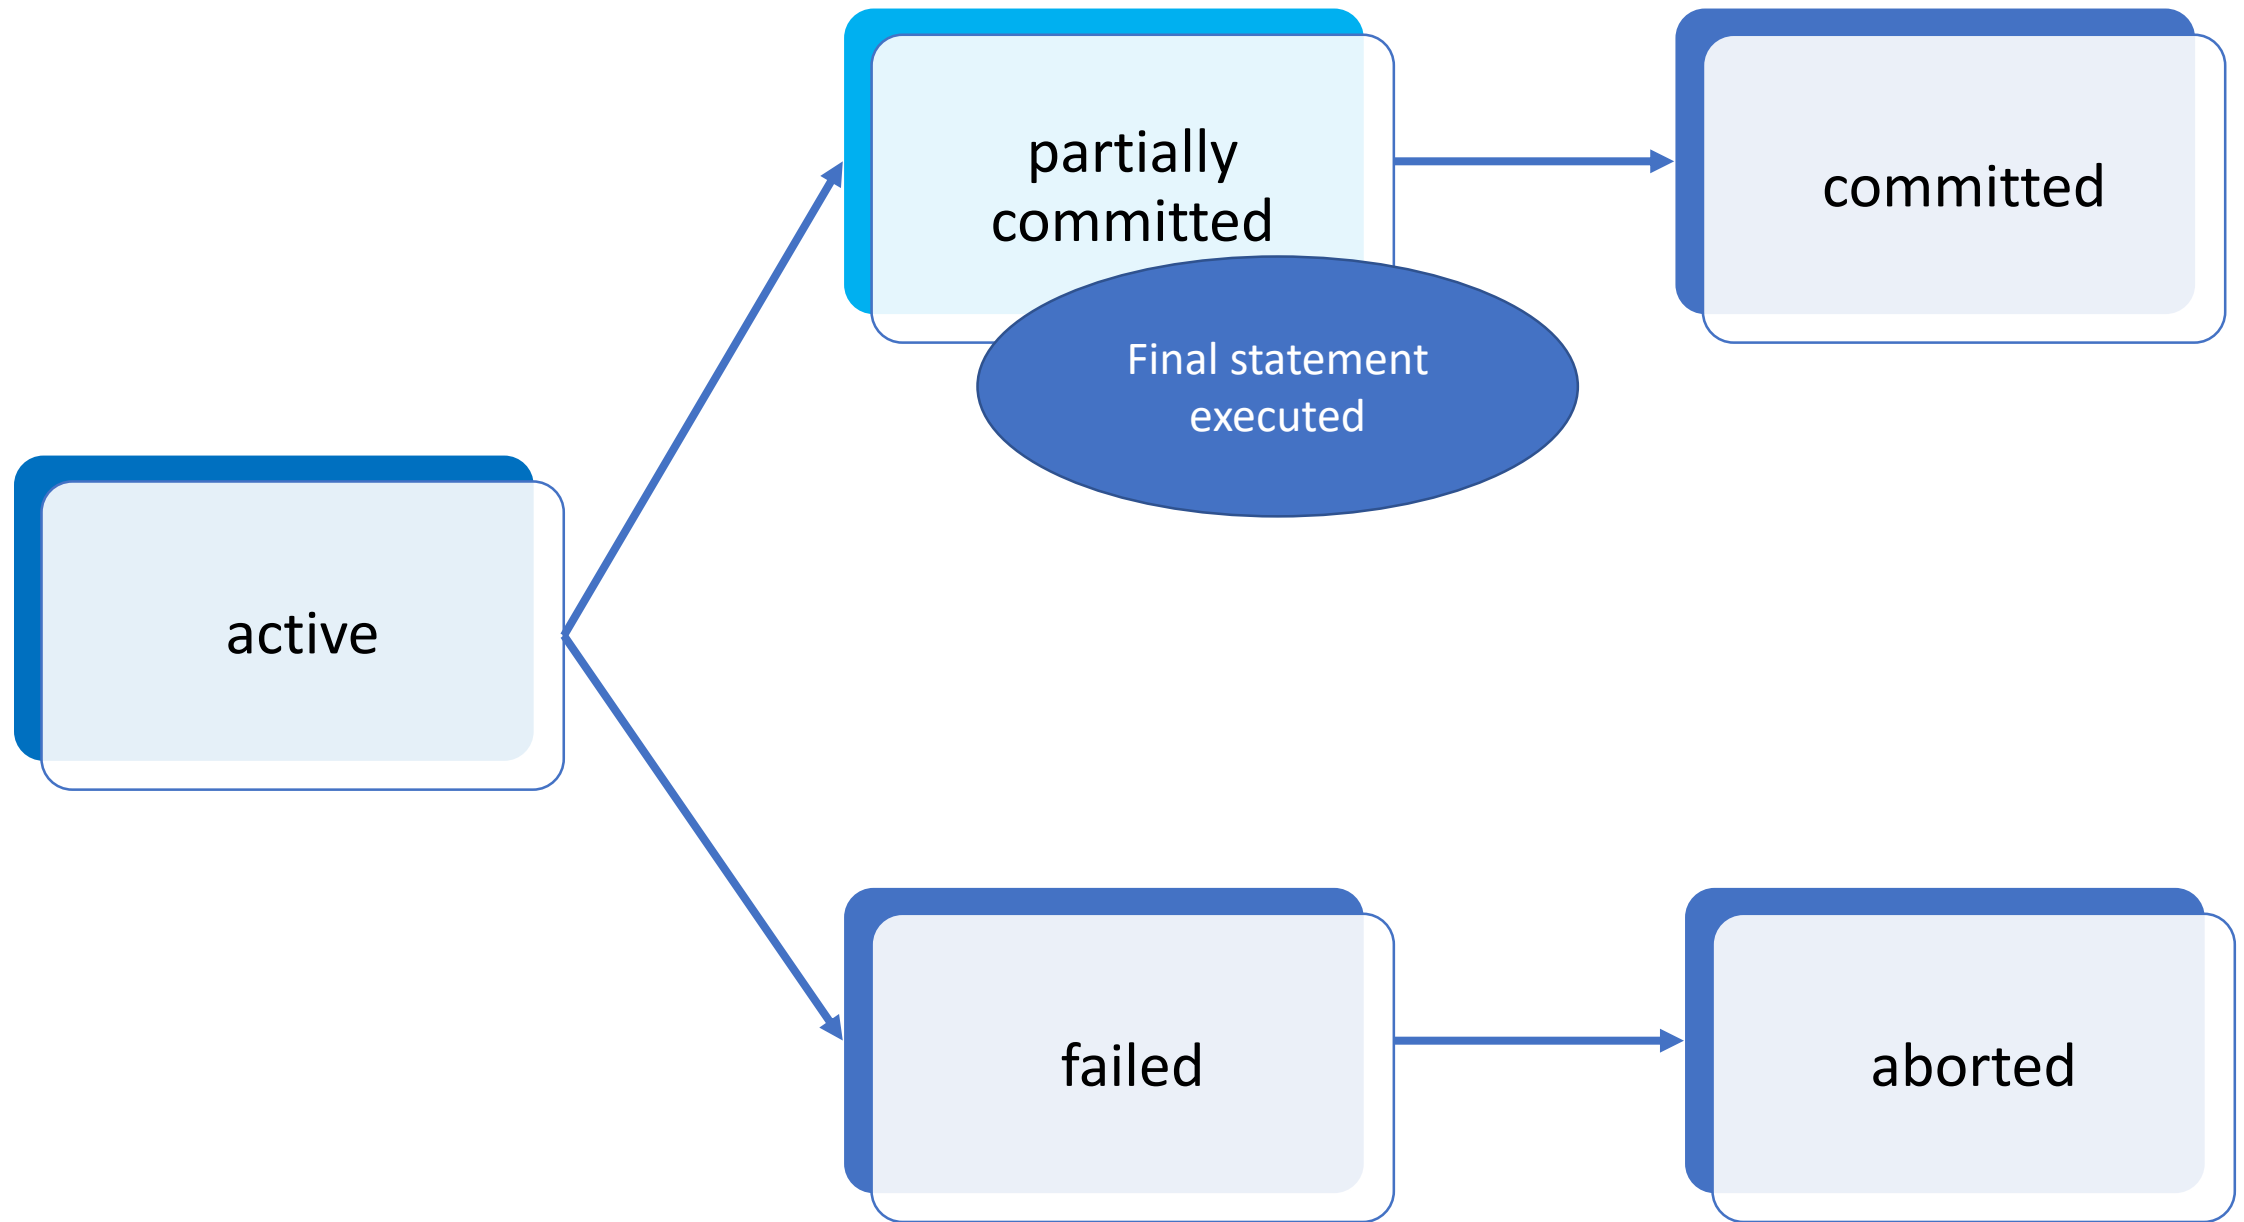
→ *recovery system*

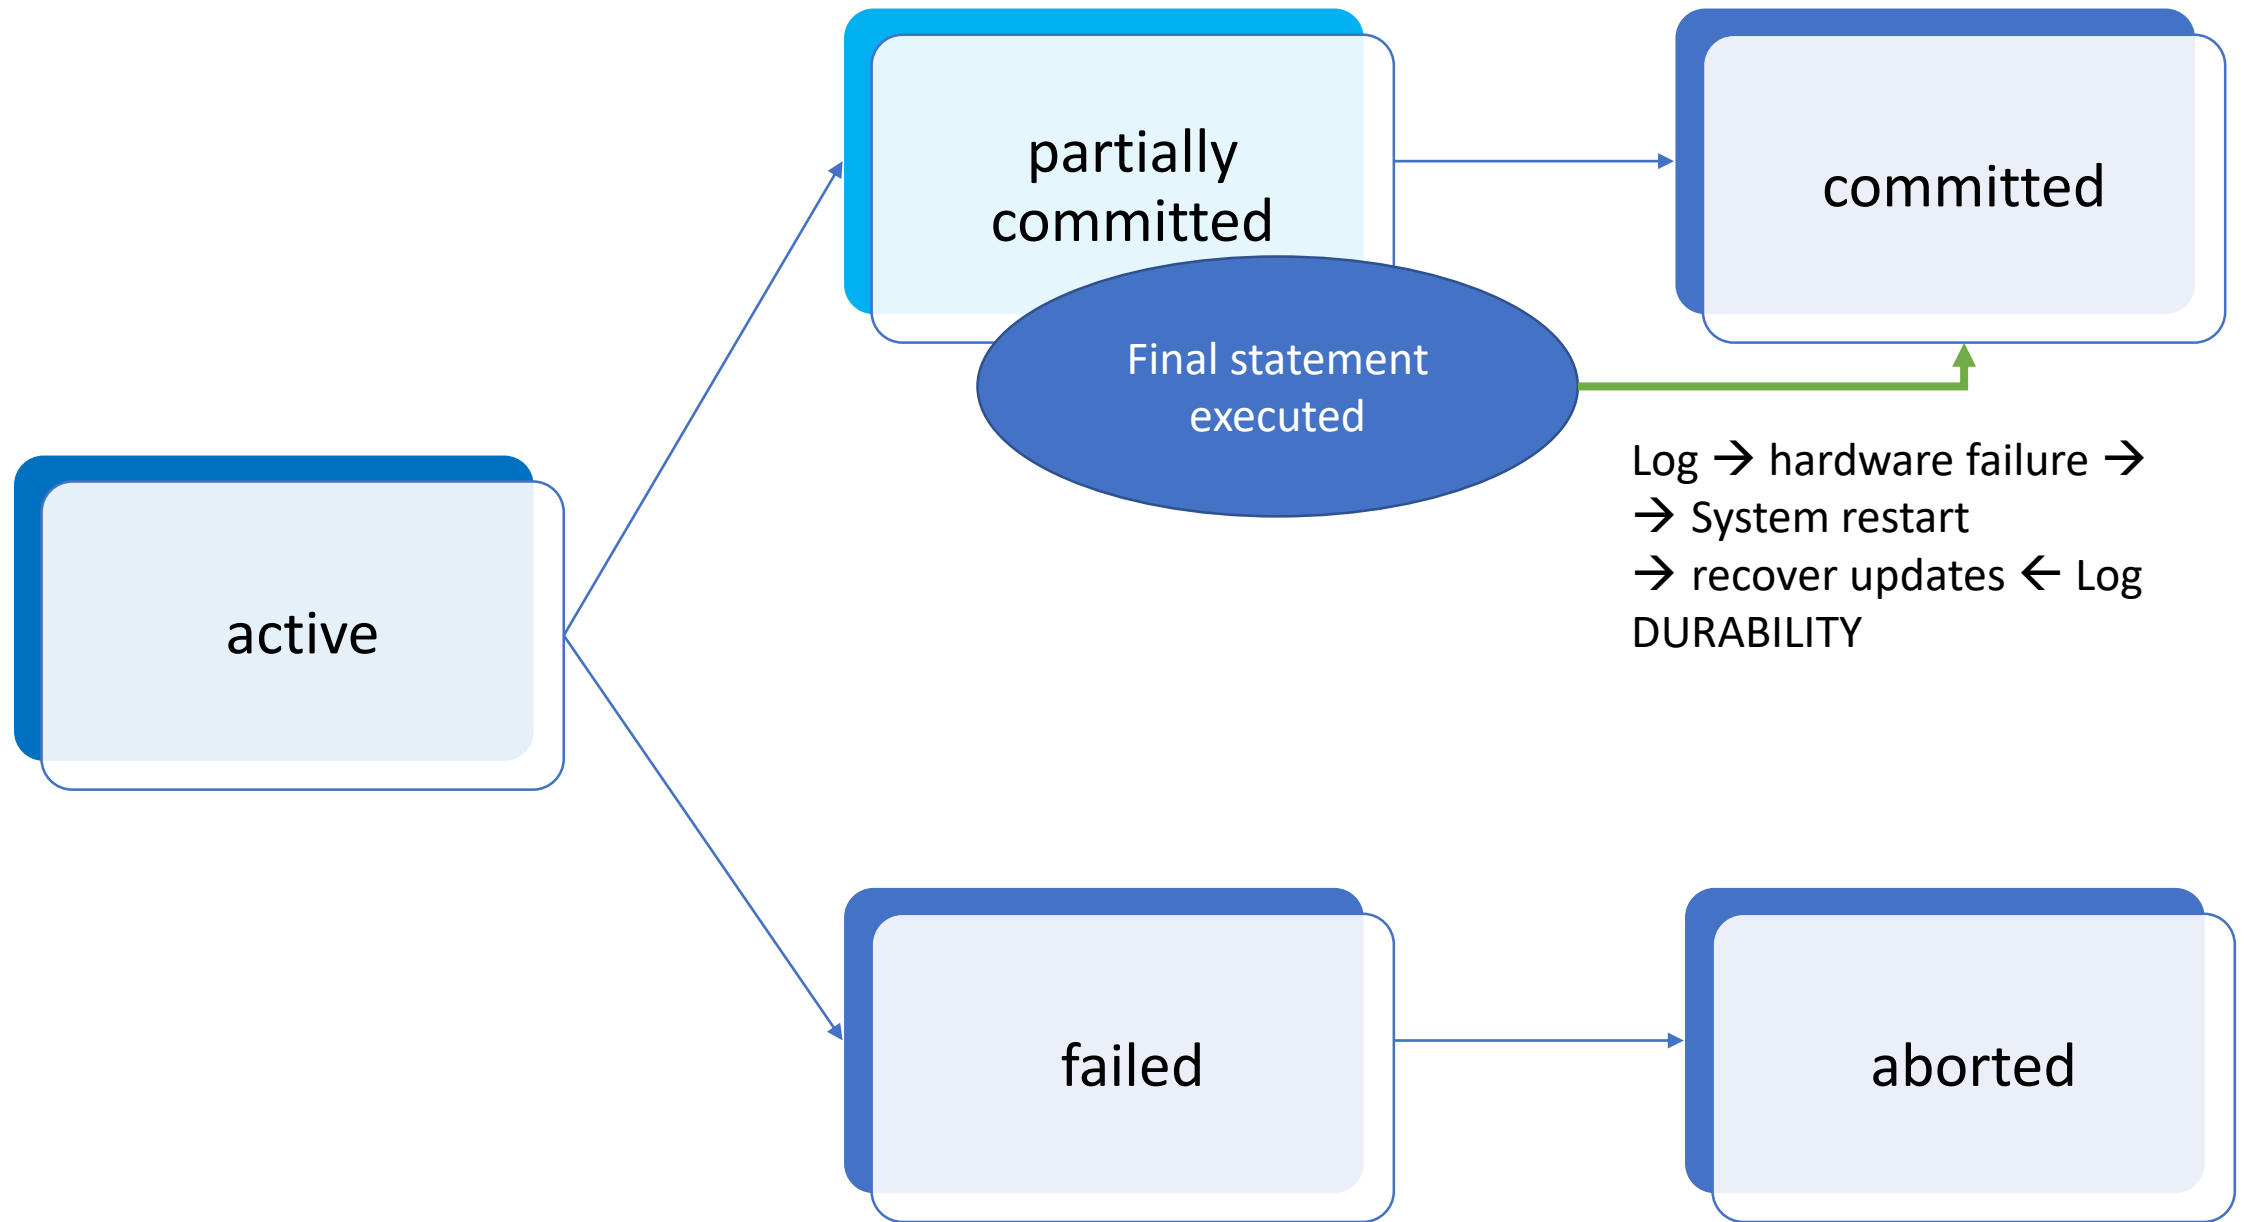
- Please answer [www.menti.com](https://www.menti.com) 13 52 85 Q1, Q2, Q3, Q4

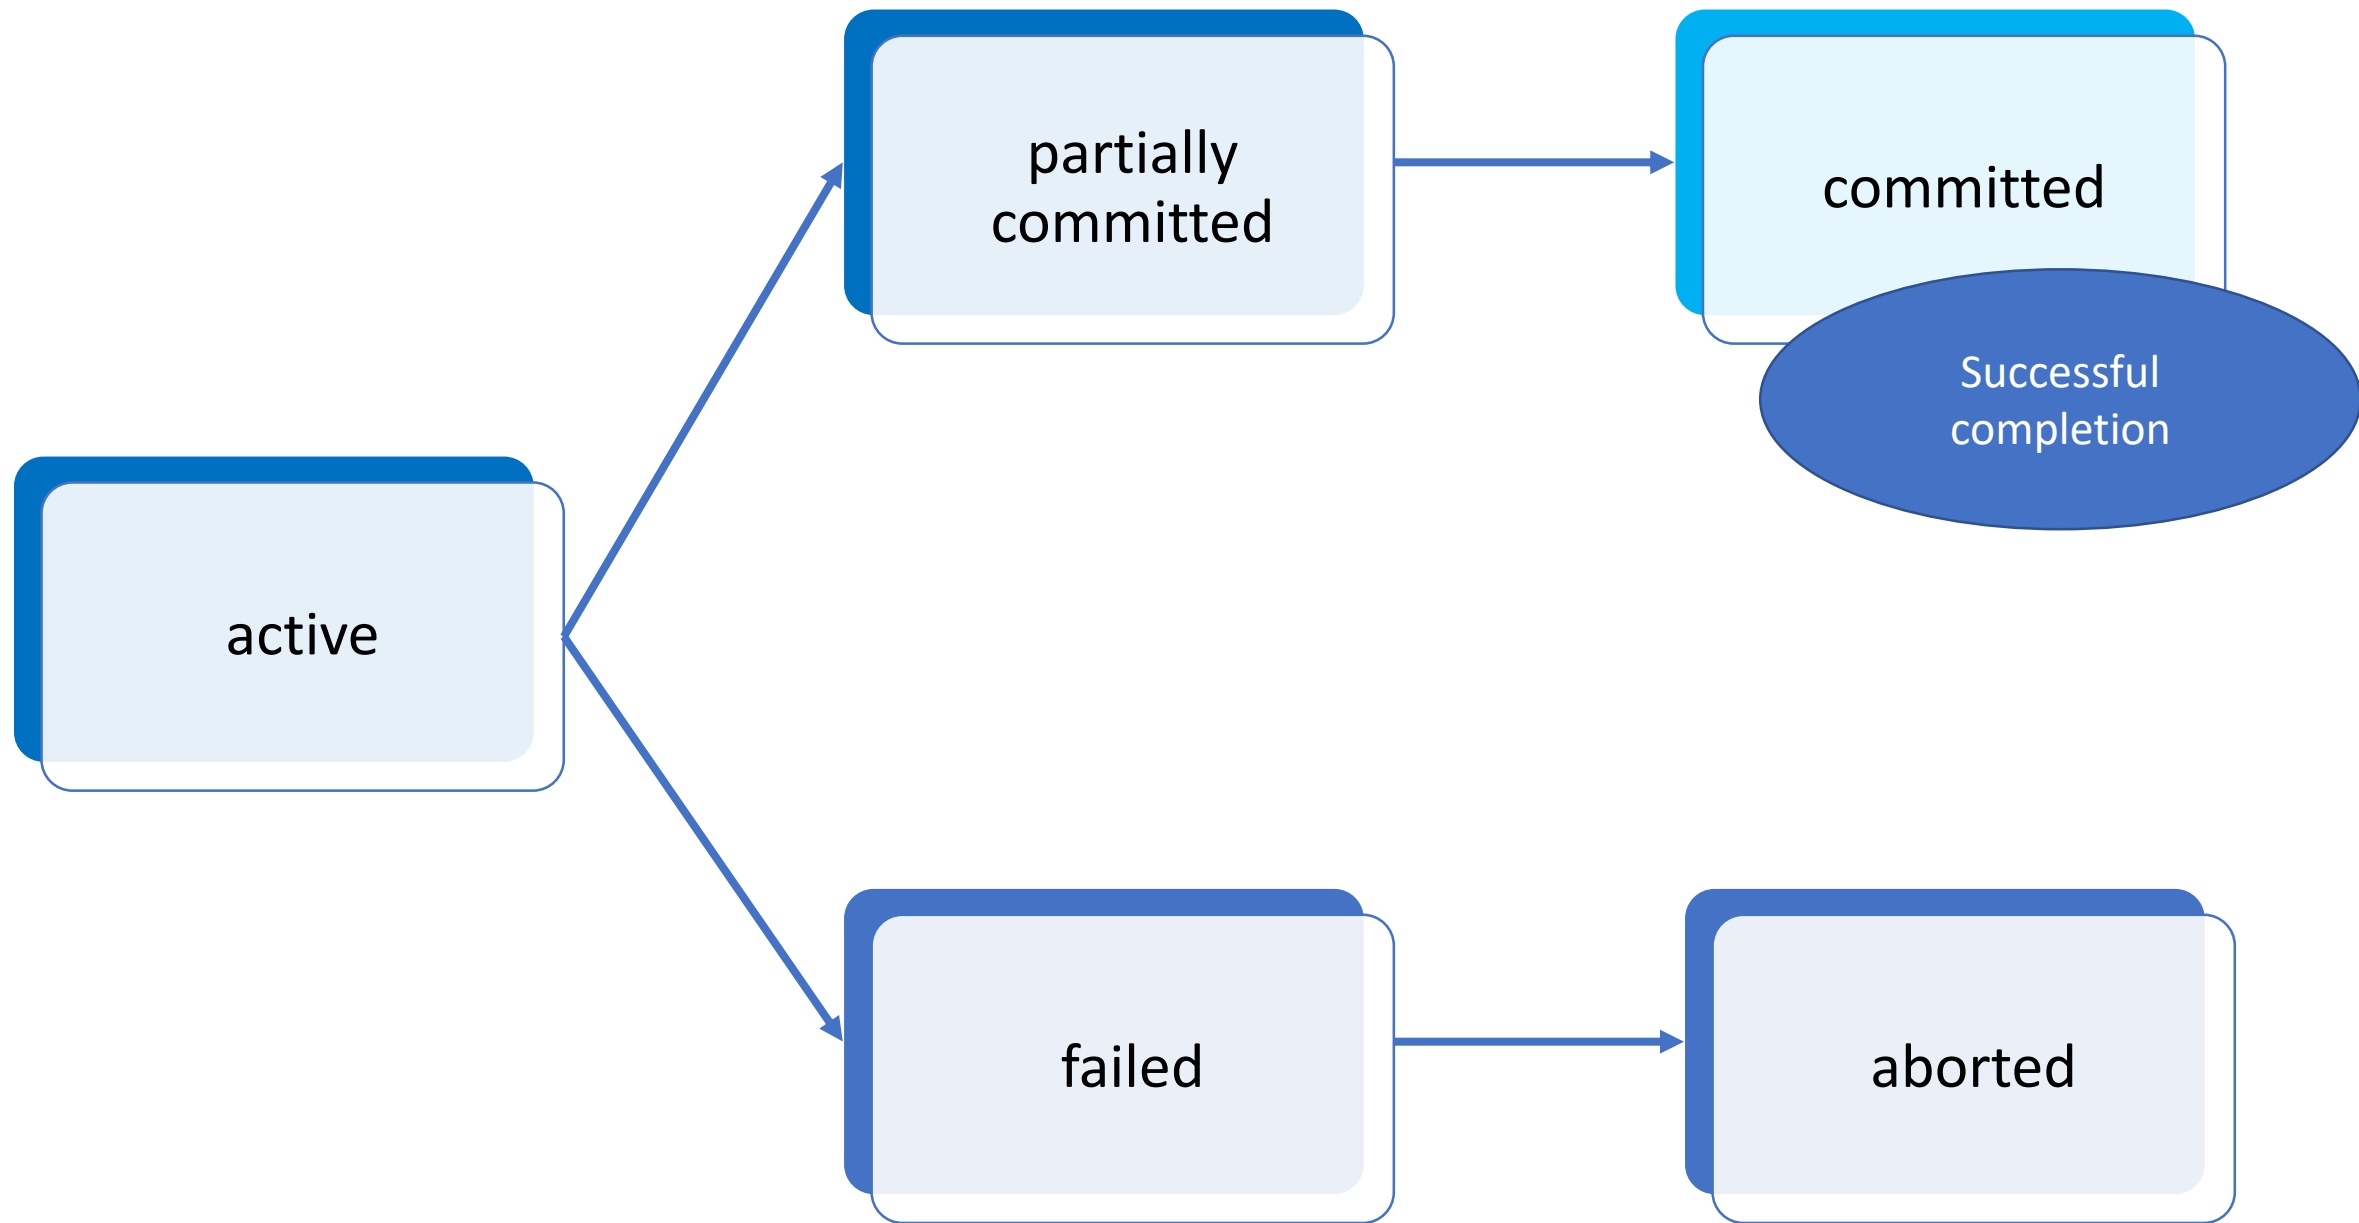
# Transaction states



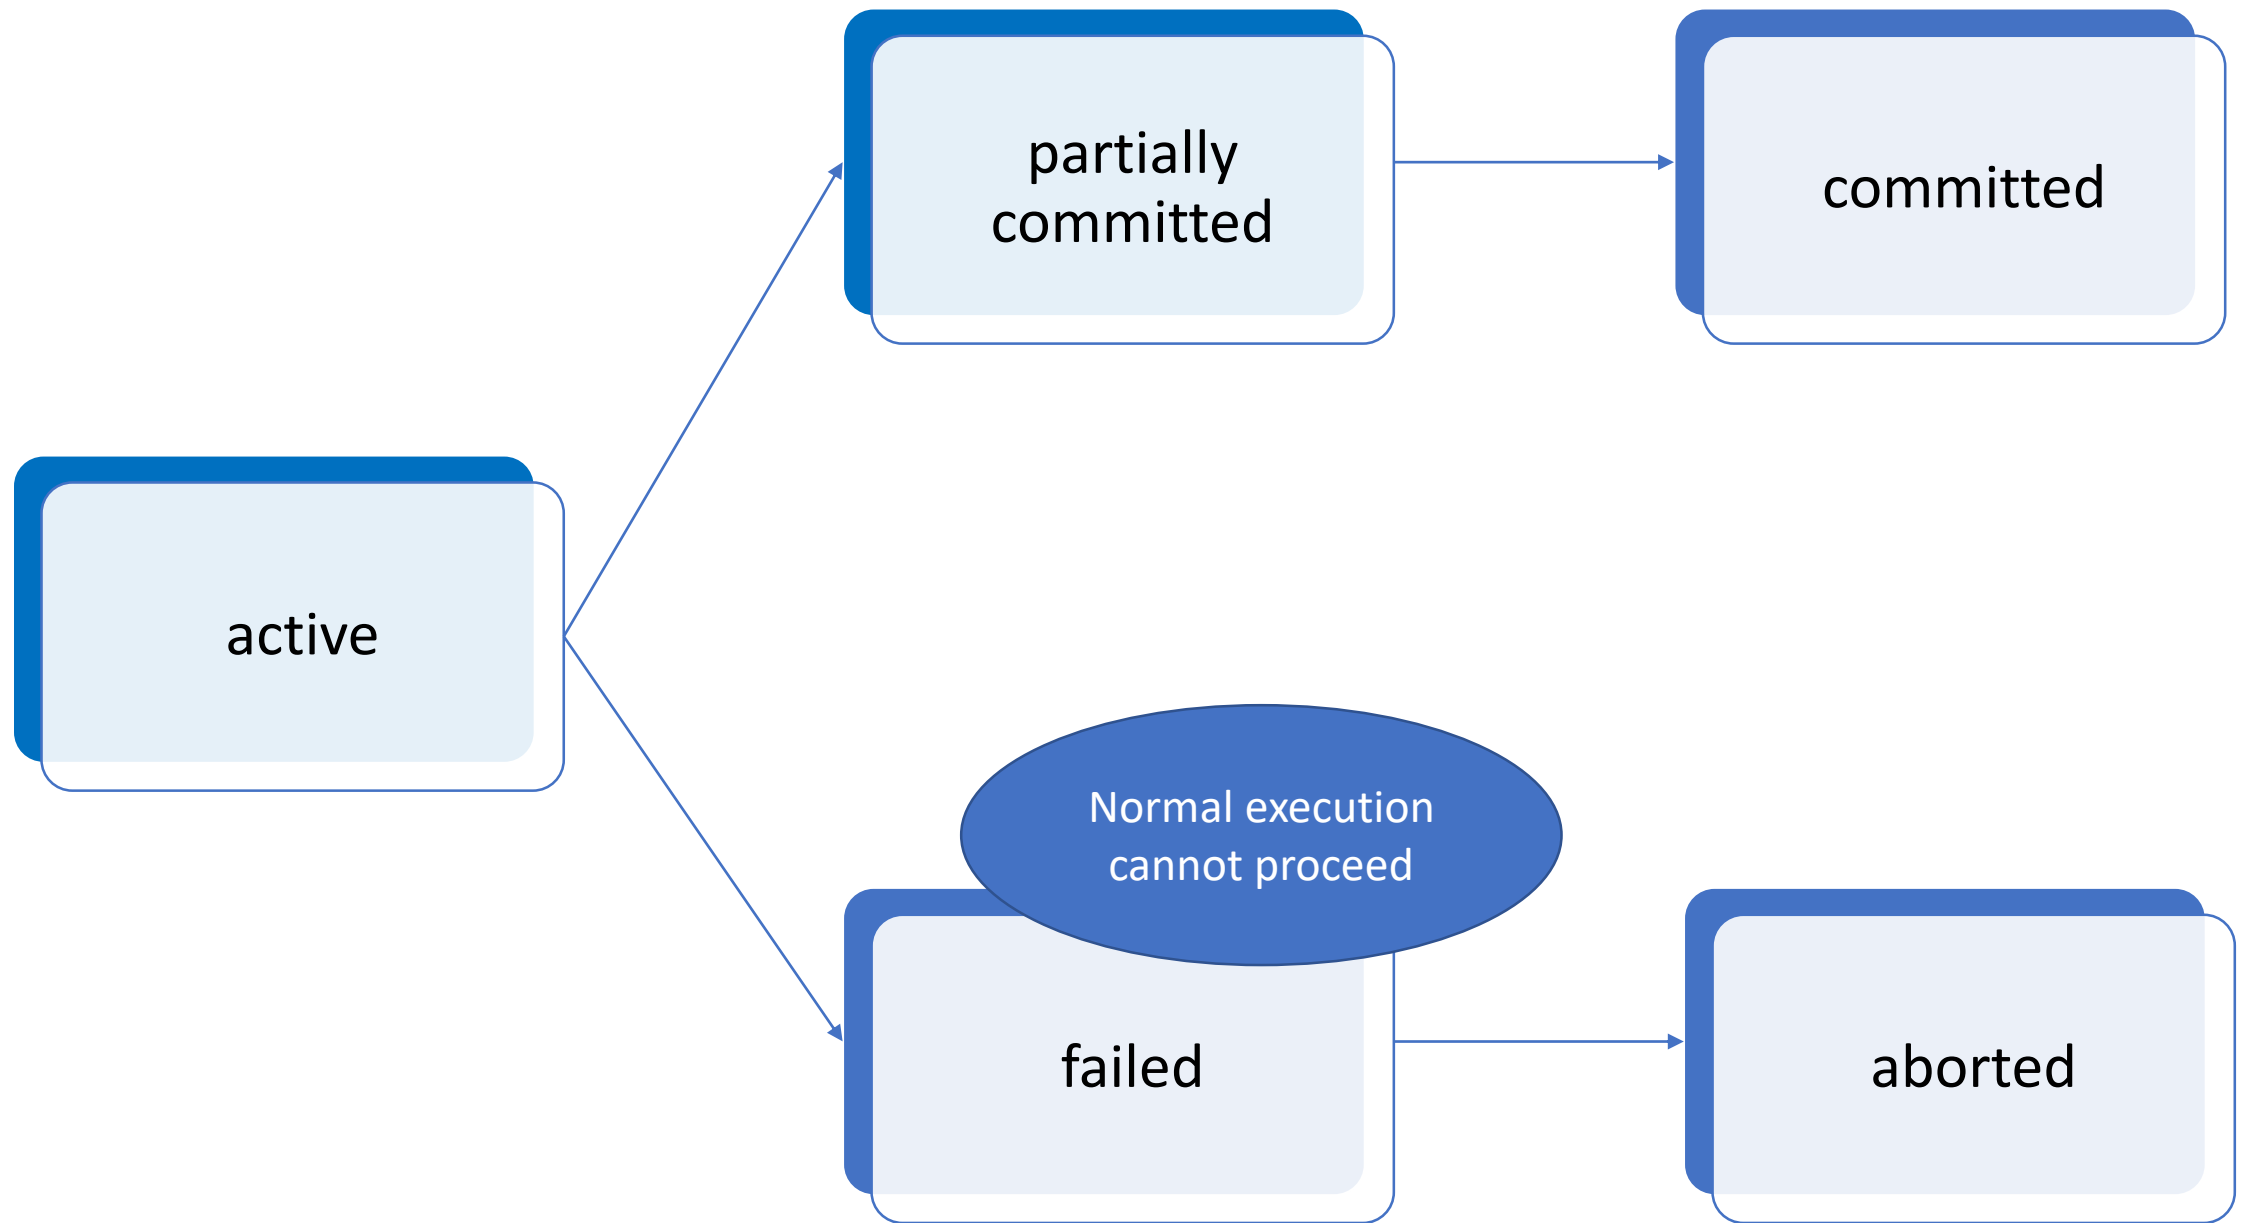


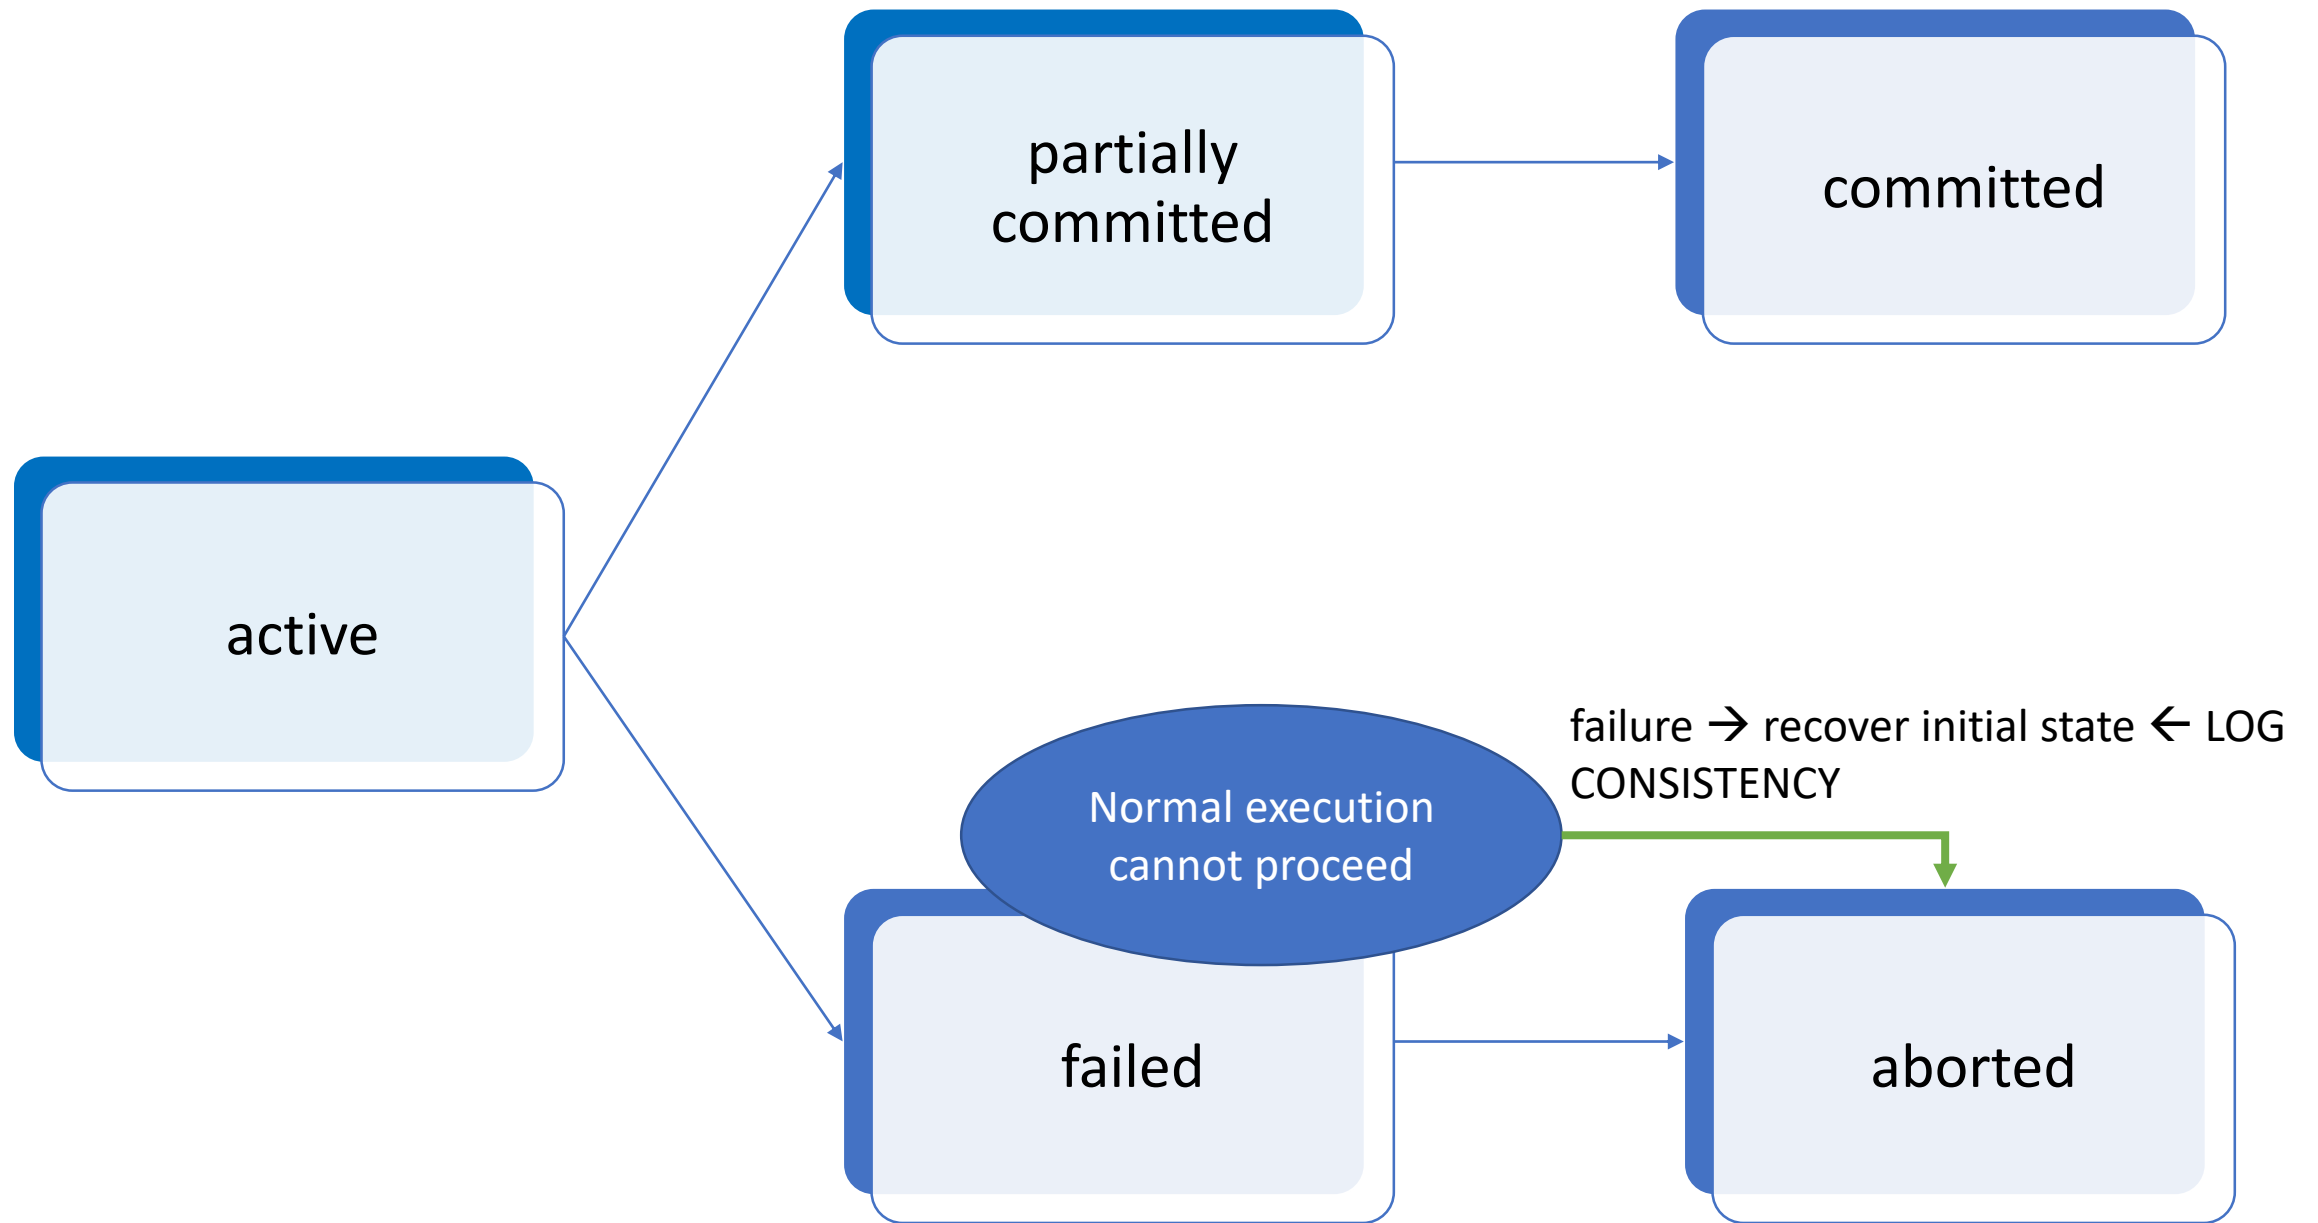


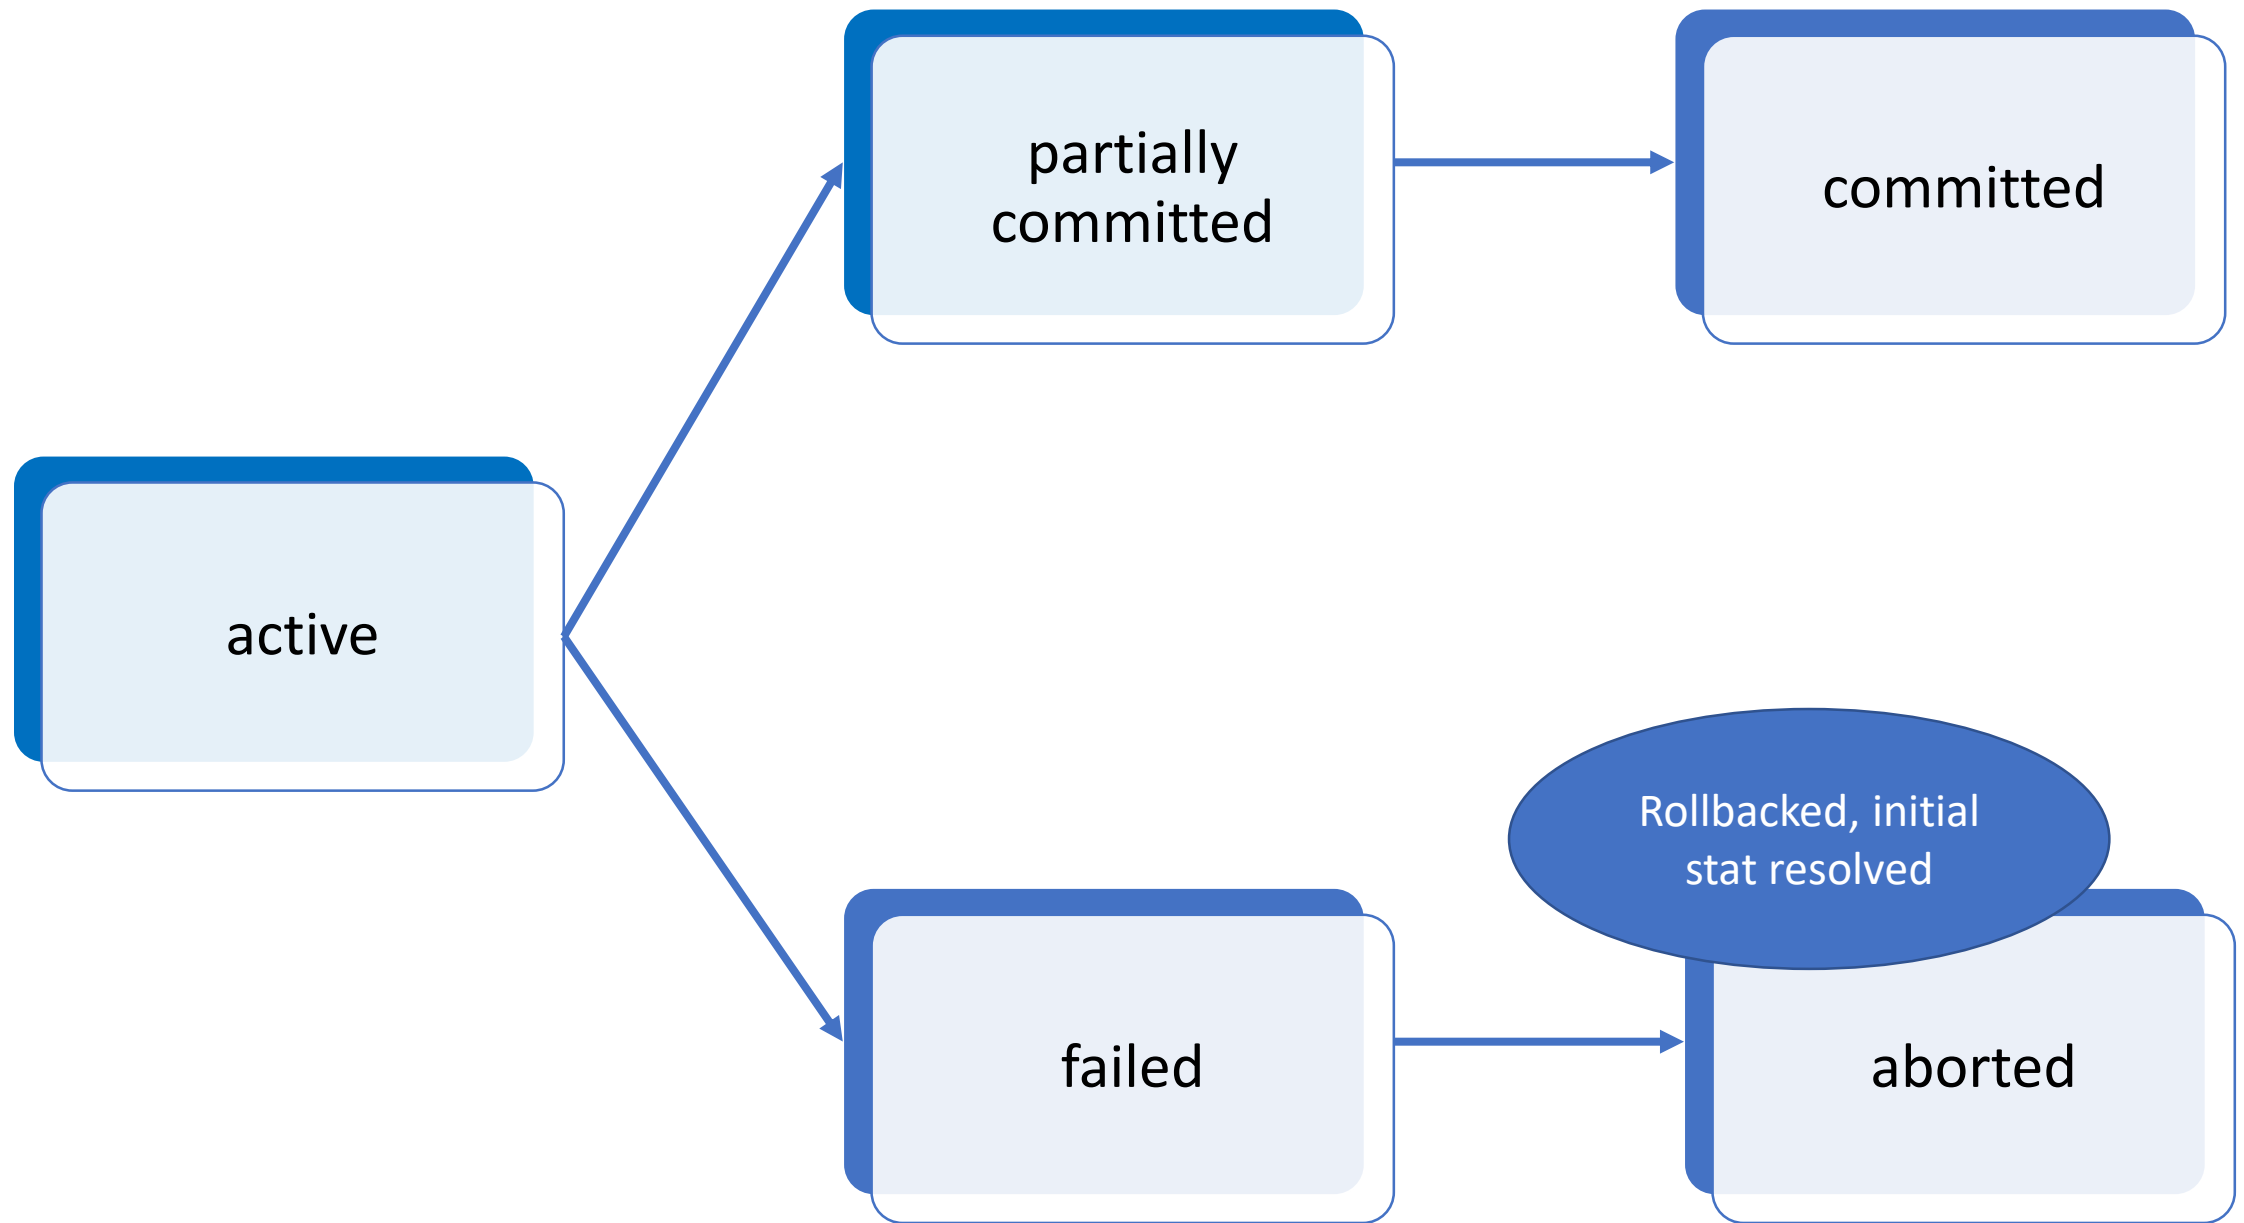












# Concurrent transactions

# Concurrent transactions

- Reduce response time: time for a transaction to be completed.
- Improved workload/resource utilization.
- ISOLATION may be violated → as a result database may be found in an inconsistent state  
→ *Concurrency control*

# Concurrent transactions - *conflicts*

- Serial execution preserves consistency, assuming that transactions preserve consistency.

first statement of transaction  $T_i$  is executed after  $T_j$  finished or

first statement of transaction  $T_j$  is executed after  $T_i$  finished

single threaded transactions

- Instructions  $I$  of  $T_i$  and  $J$  of  $T_j$  conflict  $\Leftrightarrow$  there exists a *data* accessed by both  $I$  and  $J$ , and at least one of  $I$  and  $J$  write *data*.

|                                    |                                 |                             |
|------------------------------------|---------------------------------|-----------------------------|
| 1. $I = \text{read}(\text{data})$  | $J = \text{read}(\text{data})$  | $I$ and $J$ don't conflict. |
| 2. $I = \text{read}(\text{data})$  | $J = \text{write}(\text{data})$ | conflict                    |
| 3. $I = \text{write}(\text{data})$ | $J = \text{read}(\text{data})$  | conflict                    |
| 4. $I = \text{write}(\text{data})$ | $J = \text{write}(\text{data})$ | conflict                    |

# Concurrent transactions -- Schedules

- Schedules: sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions.
  - A schedule must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instructions as the last statement
  - By default transaction assumed to execute commit instruction as its last step.
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement.

# Schedules example S1

- Serial execution.
- No conflicts.
- DB in consistent state
  - $A_{\text{new}} + B_{\text{new}} = A_{\text{old}} + B_{\text{old}}$

| T1  | T2  |
|---|---|
| <pre>read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) commit</pre> |   |
|   | <pre>read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) commit</pre> |



# Schedules example S2

- Not a serial execution.
- Equivalent to Schedule S1.
- DB in consistent state
  - $A_{\text{new}} + B_{\text{new}} = A_{\text{old}} + B_{\text{old}}$

| T1   | T2  |
|--|---|
| read (A)<br>A := A - 50<br>write (A)           |   |
|  | read (A)<br>temp := A * 0.1<br>A := A - temp<br>write (A) |
| read (B)<br>B := B + 50<br>write (B)<br>commit |   |
|  | read (B)<br>B := B + temp<br>write (B)<br>commit          |

# Concurrent transactions

- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence:
  1. Conflict serializability
  2. View serializability

# Concurrent transactions

- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence:
  1. Conflict serializability

If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are conflict equivalent.
  2. View serializability

# Schedules example S2

- Not a serial execution.
- Equivalent to Schedule S1.
- DB in consistent state
  - $A_{\text{new}} + B_{\text{new}} = A_{\text{old}} + B_{\text{old}}$

no conflict, no data item is updated by both blocks, by swapping the two blocks we obtain S1

| T1   | T2  |
|--|---|
| read (A)<br>$A := A - 50$<br>write (A)           |   |
|  | read (A)<br>$\text{temp} := A * 0.1$<br>$A := A - \text{temp}$<br>write (A) |
| read (B)<br>$B := B + 50$<br>write (B)<br>commit |   |
|  | read (B)<br>$B := B + \text{temp}$<br>write (B)<br>commit                   |

# Schedules example S3

- Not a serial execution.
- Not equivalent to Schedule S1.
- DB in inconsistent state
  - $A_{\text{new}} + B_{\text{new}} \neq A_{\text{old}} + B_{\text{old}}$

conflict, A is updated  
by both blocks

| T1  | T2  |
|---|---|
| read (A)<br>$A := A - 50$                                     |   |
|   | read (A)<br>$\text{temp} := A * 0.1$<br>$A := A - \text{temp}$<br>write (A) |
| write (A)<br>read (B)<br>$B := B + 50$<br>write (B)<br>commit |   |
|   | read (B)<br>$B := B + \text{temp}$<br>write (B)<br>commit                   |

# Concurrent transactions

1. Conflict serializability
2. View serializability

Let  $S$  and  $S'$  be 2 schedules with the same set of transactions.  $S$  and  $S'$  are view equivalent if the following 3 conditions are met, for each data item  $Q$ :

- If in schedule  $S$ , transaction  $T_i$  reads the initial value of  $Q$ , then in schedule  $S'$  also transaction  $T_i$  must read the initial value of  $Q$ .
- If in schedule  $S$  transaction  $T_i$  executes  $\text{read}(Q)$ , and that value was produced by transaction  $T_j$  (if any), then in schedule  $S'$  also transaction  $T_i$  must read the value of  $Q$  that was produced by the same  $\text{write}(Q)$  operation of transaction  $T_j$ .
- The transaction (if any) that performs the final  $\text{write}(Q)$  operation in schedule  $S$  must also perform the final  $\text{write}(Q)$  operation in schedule  $S'$ .

View equivalence is also based purely on reads and writes alone.

# Concurrent transactions

- Test serializability:

1. Conflict serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- Precedence graph — a direct graph where the vertices are the transactions (names).
- We draw an arc from  $T_i$  to  $T_j$  if the 2 transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier.
- We may label the arc by the item that was accessed.
- A schedule is CS if and only if its precedence graph is acyclic.
- If precedence graph is acyclic, the serializability order can be obtained by a **topological sorting** of the graph.

# Concurrent transactions

- Test serializability :

2. View serializability

- The problem of checking if a schedule is view serializable falls in the class of NP-complete problems. Thus, existence of an efficient algorithm is extremely unlikely.
- Practical algorithms that just check some sufficient conditions for view serializability can still be used.

Please answer [www.menti.com](https://www.menti.com) 13 52 85 Q5



# Isolation levels

# Isolation levels

- **Isolation:** execute a transaction *as if* there are no other concurrent transactions running simultaneously.
  - Prevent read or write of incorrect, temporary, aborted data processed by concurrent transactions
- **Isolation levels:** trade off between *perfect* isolation and performance
  - response time: time before a transaction completes
  - throughput: number of transactions per second

# Level **Serializability**, perfect isolation

- The final state of the database is equivalent to a state of the database if the transactions were run sequentially.
  - serializable schedule
- Way of obtaining serializability:
  - locking
  - timestamp validation
  - multi-versioning

# Transactions errors

lost-update anomaly

final stock 12!

| T1   | T2   |
|--|--|
| <pre>select qte into :nS from stock where n_prod = 100 --nS = 13</pre> |  |
|  | <pre>select qte into :nS from stock where n_prod = 100</pre> |
| <pre>update stock set qte = :nS - 1 where n_prod = 100</pre>           |  |
|  | <pre>update stock set qte = :nS - 1 where n_prod = 100</pre> |
| <pre>insert into orders(n_prod, qte) values(100, 1) commit</pre>       |  |
|  | <pre>insert into orders(n_prod, qte) values(100, 1)</pre>    |

# Transactions errors

**dirty-read anomaly**  
number of products  
ordered + qte\_stock !=  
initial stock  
1 product missing!  
Read uncommitted  
data

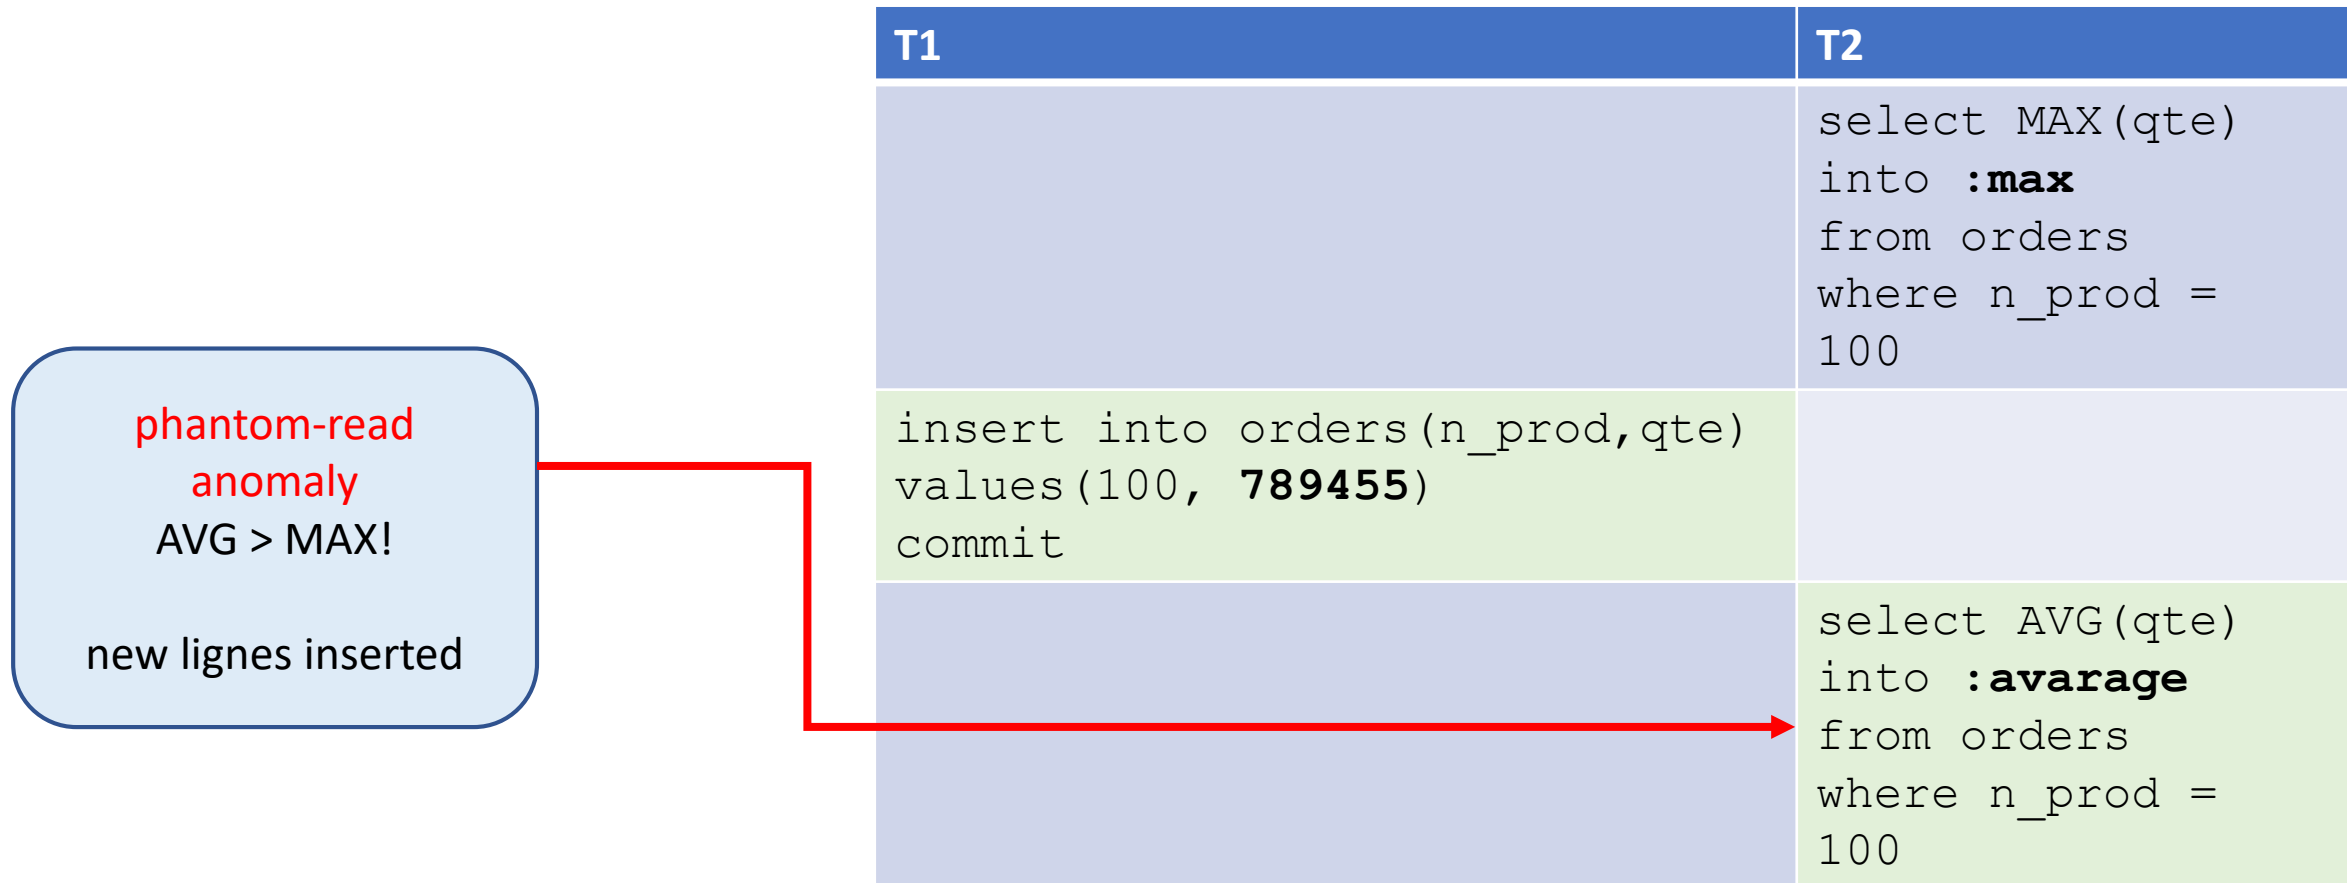
| T1   | T2   |
|--|--|
| <pre>select qte into :nS from stock where n_prod = 100 --nS = 13</pre> |  |
| <pre>update stock set qte = :nS - 1 where n_prod = 100</pre>           |  |
|  | <pre>select sum(qte) into :nO from orders where n_prod = 100</pre>               |
|  | <pre>select qte into :nS from stock where n_prod = 100 --nO+nS!=init_stock</pre> |
| <pre>insert into orders(n_prod, qte) values(100, 1) commit</pre>       |  |

# Transactions errors

**non-repeatable read  
anomaly**  
only one insert into  
restock is needed!  
read twice, different  
values

| T1  | T2   |
|---|--|
| ...   | <pre>select qte into :nS from stock where n_prod = 100 --nS = 10</pre>   |
|   | <pre>if nS &lt; 15 and nS &gt;= 10   insert into restock(n_prod, qte)   values(100, 5)</pre>                                     |
| <pre>update stock set qte = :nS - 1 where n_prod = 100  insert into orders (n_prod, qte) values(100, 1)  commit</pre> |  |
|   | <pre>select qte into :nS from stock where n_prod = 100  if nS &lt; 10   insert into restock(n_prod, qte)   values(100, 15)</pre> |

# Transactions errors



# Transactions errors

**dirty-write anomaly**  
final stock 11! In the first transaction, the stock returns to 13. Only one update should decrease the number of products.

| T1   | T2   |
|--|--|
| <pre>select qte into :nS from stock where n_prod = 100 --nS = 13</pre> |  |
| <pre>update stock set qte = :nS - 1 where n_prod = 100</pre>           |  |
|  | <pre>select qte into :nS from stock where n_prod = 100 --nS = 12</pre> |
|  | <pre>update stock set qte = :nS - 1 where n_prod = 100 --nS = 11</pre> |
| abort  |  |
|  | <pre>insert ... commit</pre>   |







# Isolation levels

- weaker the isolation level → more anomalies may occur





|                     | ERROR | lost-update   | dirty-reads   | non-repeatable<br>reads   | phantom   |
|---------------------|-------|---|---|---|---|
| LEVEL               |       |   |   |   |   |
| READ<br>UNCOMMITTED |       |    |    |    |    |
| READ<br>COMMITTED   |       |    |    |    |    |
| REPEATABLE<br>READ  |       |   |   |   |   |
| SERIALIZABLE        |       |  |  |  |  |

# Isolation levels

|                    | ERROR | lost-update   | dirty-reads   | non-repeatable<br>reads   | phantom   |
|--------------------|-------|---|---|---|---|
| LEVEL              |       |   |   |   |   |
| REPEATABLE<br>READ |       |  |  |  |  |





- read only committed
- between two reads of an item by a transaction, no other transaction is allowed to update it.
- a transaction may find other data inserted by a committed transaction

# Isolation levels

|                   | ERROR | lost-update   | dirty-reads   | non-repeatable<br>reads   | phantom   |
|-------------------|-------|---|---|---|---|
| LEVEL             |       |   |   |   |   |
| READ<br>COMMITTED |       |  |  |  |  |

- read only committed
- does not require repeatable reads. Between two reads of a data item by the transaction, another transaction may have updated the data item and committed.

# Isolation levels

|                     | ERROR | lost-update   | dirty-reads   | non-repeatable<br>reads   | phantom   |
|---------------------|-------|---|---|---|---|
| LEVEL               |       |   |   |   |   |
| READ<br>UNCOMMITTED |       |  |  |  |  |

- allows uncommitted data to be read
- all the isolation levels prevent writes to a data item that has already been written by another transaction not yet committed or aborted (rollback).
- Please answer [www.menti.com](https://www.menti.com) 13 52 85 Q6, Q7

# Achieving isolation

- Versioning
  - Transactions read from a “snapshot” of the database.
- Locking
- Timestamp

# Locking

- Locks prevent destructive interactions between transactions accessing the same resource.
  - Shared      access to read
  - Exclusive    access to read and write
- Locks (Shared, Shared) compatible.
- Locks (Shared, Exclusive) not compatible.
- A transaction waits until all incompatible locks held by other transactions are released.
- <https://oracle-base.com/articles/misc/deadlocks>
- [https://docs.oracle.com/cd/B19306\\_01/server.102/b14220/consist.htm](https://docs.oracle.com/cd/B19306_01/server.102/b14220/consist.htm)

# Snapshot isolation

- Snapshot of the database at the beginning of each transaction.
- The transaction operates only on that snapshot.
- The snapshot consists only of committed values.
- Updates are kept in transaction workspace until commit.
- Implemented with timestamp-versioning

# Consistency levels

To be added, more info in the following video.



# BASE

NoSql consistency model

To be added, more info in the following video.